

Create a Kubernetes cluster using minikube

Install Minikube on Mac CLI

To install minikube on x86–64 mac OS, run the following two commands:

```
curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-darwin-amd64  
sudo install minikube-darwin-amd64 /usr/local/bin/minikube
```

Use the following command to start a cluster:

```
minikube start
```

Use the following command to confirm minikube is running:

```
minikube status
```

```
[Brandis-MBP:W17complex brandi$ minikube status  
minikube  
type: Control Plane  
host: Running  
kubelet: Running  
apiserver: Running  
kubeconfig: Configured
```

Create a 3 Node Cluster

We want to create a 3 node cluster that will consist of one master node (for the rest of the tutorial we will call this node “control-plane”) and two worker nodes. We can create a 3 node cluster by using the following command:

```
minikube start --nodes 3 -p <cluster_name>
```

```

Use "minikube options" for a list of global command-line options (applies to all commands).
Brandis-MBP:W17complex brandi$ minikube start --nodes 3 -p k8cluster
[k8cluster] minikube v1.29.0 on Darwin 12.6.3
🌟 Automatically selected the docker driver
👍 Using Docker Desktop driver with root privileges
👍 Starting control plane node k8cluster in cluster k8cluster
🔄 Pulling base image ...
🔥 Creating docker container (CPUs=2, Memory=2200MB) ...
🌐 Preparing Kubernetes v1.26.1 on Docker 20.10.23 ...
  ▪ Generating certificates and keys ...
  ▪ Booting up control plane ...
  ▪ Configuring RBAC rules ...
🔗 Configuring CNI (Container Networking Interface) ...
  ▪ Using image gcr.io/k8s-minikube/storage-provisioner:v5
🔍 Verifying Kubernetes components...
🌟 Enabled addons: storage-provisioner, default-storageclass

👍 Starting worker node k8cluster-m02 in cluster k8cluster
🔄 Pulling base image ...
🔥 Creating docker container (CPUs=2, Memory=2200MB) ...
🌐 Found network options:
  ▪ NO_PROXY=192.168.94.2
  ⚠ This container is having trouble accessing https://registry.k8s.io
  To pull new external images, you may need to configure a proxy: https://minikube.sigs.k8s.io/docs/reference/networking/proxy/
🌐 Preparing Kubernetes v1.26.1 on Docker 20.10.23 ...
  ▪ env NO_PROXY=192.168.94.2
🔍 Verifying Kubernetes components...

👍 Starting worker node k8cluster-m03 in cluster k8cluster
🔄 Pulling base image ...
🔥 Creating docker container (CPUs=2, Memory=2200MB) ...
🌐 Found network options:
  ▪ NO_PROXY=192.168.94.2,192.168.94.3
🌐 Preparing Kubernetes v1.26.1 on Docker 20.10.23 ...
  ▪ env NO_PROXY=192.168.94.2
  ▪ env NO_PROXY=192.168.94.2,192.168.94.3
🔍 Verifying Kubernetes components...
👍 Done! kubectl is now configured to use "k8cluster" cluster and "default" namespace by default

```

The cluster may take several minutes to create, but minikube will output updates as things are happening. Once it says “Done!” use the following command to see the three nodes you just created.

```
kubectl get nodes
```

```

Brandis-MBP:W17complex brandi$ kubectl get nodes
NAME                STATUS    ROLES                  AGE     VERSION
k8cluster            Ready     control-plane          14m     v1.26.1
k8cluster-m02        Ready     <none>                  12m     v1.26.1
k8cluster-m03        Ready     <none>                  10m     v1.26.1

```

3 node cluster

Label Nodes

When we deploy our Redis and Apache pods, we do not want them to deploy to our control-plane, so we need to label our second and third nodes as “worker”. Use the following command to apply a worker label to the k8cluster-m02 and k8cluster-m03 nodes. You will have to run the command twice, once for each of the above nodes.

```
kubectl label node <node_name> node-role.kubernetes.io/worker=worker
```

```
Brandis-MBP:W17complex brandi$ kubectl label node k8cluster-m02 node-role.kubernetes.io/worker=worker
node/k8cluster-m02 labeled
Brandis-MBP:W17complex brandi$ kubectl label node k8cluster-m03 node-role.kubernetes.io/worker=worker
node/k8cluster-m03 labeled
```

Use the following command to view the newly labeled nodes.

```
kubectl get nodes
```

```
Brandis-MBP:W17complex brandi$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8cluster	Ready	control-plane	18m	v1.26.1
k8cluster-m02	Ready	worker	15m	v1.26.1
k8cluster-m03	Ready	worker	13m	v1.26.1

The YAML files that we will deploy will search for these worker nodes based on a key:value label pair. Use the following command to apply a key:value label to the worker nodes:

```
kubectl label nodes <node_name> role=worker
```

```
[Brandis-MBP:W17complex brandi$ kubectl label nodes k8cluster-m02 role=worker
node/k8cluster-m02 labeled
```

Run the command again for the second worker node (k8cluster-m03).

```
[Brandis-MBP:W17complex brandi$ kubectl label nodes k8cluster-m03 role=worker  
node/k8cluster-m03 labeled
```

Redis Deployment YAML File

In my last tutorial [Create a Kubernetes Cluster Using Docker Desktop](#), I went into detail on how to create a Kubernetes deployment YAML file, If you need to create a YAML file for the Redis and Apache deployments, please visit that tutorial now and come back when your files are created. Alternatively, you can just copy the code below for the Redis deployment with four replicas and save it as a .yaml file.

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: redis-deploy  
  labels:  
    app: redis  
spec:  
  replicas: 4  
  selector:  
    matchLabels:  
      app: redis  
  template:  
    metadata:  
      labels:  
        app: redis  
    spec:  
      containers:  
        - name: redis  
          image: redis:7.0.9  
          ports:  
            - containerPort: 6379  
      nodeSelector:  
        role: worker
```

Notice in the file we've added **nodeSelector** on the container being created. With `nodeSelector`, we can choose which node our containers will be deployed on. Earlier we assigned our second and third nodes to be "worker" nodes and indicated that with the key:value pair `role:worker`. When we deploy this YAML file, the engine will look for any nodes that have the `role:worker` label and will deploy replicas on those nodes only. By using these labels, we are avoiding deploying replicas on the control-panel.

Apache Deployment YAML File

Here is the file contents for the Apache deployment with ten replicas.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: httpd-deploy
  labels:
    app: httpd
spec:
  replicas: 10
  selector:
    matchLabels:
      app: httpd
  template:
    metadata:
      labels:
        app: httpd
    spec:
      containers:
        - name: httpd
          image: httpd:2.4.56
          ports:
            - containerPort: 80
          nodeSelector:
            role: worker
```

Note that these two files can actually be combined into one file by using `---` to separate each manifest. To see what a single file would look like, visit my GitHub [here](#).

Deploy Redis and Apache Pods

Using the CLI, navigate to the directory that contains your Redis and Apache YAML files. Use the following command to deploy the pods:

```
kubectl apply -f <yaml_file.yml>
```

You will need to run this command twice. Once for the Redis YAML file and once for the Apache YAML file.

```
Brandis-MBP:W17complex brandi$ kubectl apply -f k8_httpd.yml
deployment.apps/httpd-deploy created
Brandis-MBP:W17complex brandi$ kubectl apply -f k8s_redis.yml
deployment.apps/redis-deploy created
```

Your pods containing your replicas should have successfully deployed. To confirm this, use the following command:

```
kubectl get pods -o wide
```

```
Brandis-MBP:W17complex brandi$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
httpd-deploy-6765c96878-4wbb7	1/1	Running	0	2m45s	10.244.2.4	k8cluster-m03	<none>		<none>	
httpd-deploy-6765c96878-5fjnb	1/1	Running	0	2m46s	10.244.2.6	k8cluster-m03	<none>		<none>	
httpd-deploy-6765c96878-5gjmj	1/1	Running	0	2m45s	10.244.1.8	k8cluster-m02	<none>		<none>	
httpd-deploy-6765c96878-6d4m6	1/1	Running	0	2m45s	10.244.2.8	k8cluster-m03	<none>		<none>	
httpd-deploy-6765c96878-7cdkl	1/1	Running	0	2m45s	10.244.1.7	k8cluster-m02	<none>		<none>	
httpd-deploy-6765c96878-896qx	1/1	Running	0	2m45s	10.244.2.7	k8cluster-m03	<none>		<none>	
httpd-deploy-6765c96878-cfrs6	1/1	Running	0	2m45s	10.244.1.6	k8cluster-m02	<none>		<none>	
httpd-deploy-6765c96878-1fhsr	1/1	Running	0	2m45s	10.244.1.5	k8cluster-m02	<none>		<none>	
httpd-deploy-6765c96878-q4n7k	1/1	Running	0	2m45s	10.244.2.5	k8cluster-m03	<none>		<none>	
httpd-deploy-6765c96878-shjpd	1/1	Running	0	2m45s	10.244.1.4	k8cluster-m02	<none>		<none>	
redis-deploy-6ff456db6f-98jgg	1/1	Running	0	3m48s	10.244.2.3	k8cluster-m03	<none>		<none>	
redis-deploy-6ff456db6f-d4dmj	1/1	Running	0	3m48s	10.244.1.3	k8cluster-m02	<none>		<none>	
redis-deploy-6ff456db6f-fptv2	1/1	Running	0	3m48s	10.244.1.2	k8cluster-m02	<none>		<none>	
redis-deploy-6ff456db6f-gbcdb	1/1	Running	0	3m48s	10.244.2.2	k8cluster-m03	<none>		<none>	

You can see that all fourteen replicas are running and they are only running on the worker nodes (k8cluster-m02 and k8cluster-m03).

Get Brandi McCall's stories in your inbox

Join Medium for free to get updates from this writer.

Subscribe

To get more information about what we just created, use the following command:

```
kubectl get all
```



```
[Brandis-MBP:W17complex brandi$ kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/httpd-deploy-6765c96878-4wbb7	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-5fjnb	1/1	Running	0	4m55s
pod/httpd-deploy-6765c96878-5gjmg	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-6d4m6	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-7cdk1	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-896qx	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-cfrs6	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-lfhsr	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-q4n7k	1/1	Running	0	4m54s
pod/httpd-deploy-6765c96878-shjpd	1/1	Running	0	4m54s
pod/redis-deploy-6ff456db6f-98jgg	1/1	Running	0	5m57s
pod/redis-deploy-6ff456db6f-d4dmj	1/1	Running	0	5m57s
pod/redis-deploy-6ff456db6f-fptv2	1/1	Running	0	5m57s
pod/redis-deploy-6ff456db6f-gbcdb	1/1	Running	0	5m57s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	35m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/httpd-deploy	10/10	10	10	4m55s
deployment.apps/redis-deploy	4/4	4	4	5m57s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/httpd-deploy-6765c96878	10	10	10	4m55s
replicaset.apps/redis-deploy-6ff456db6f	4	4	4	5m57s

```
[Brandis-MBP:W17complex brandi$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8cluster	Ready	control-plane	36m	v1.26.1
k8cluster-m02	Ready	worker	34m	v1.26.1
k8cluster-m03	Ready	worker	32m	v1.26.1

With this command you can view all your replicas, your cluster IP that was automatically created, your two deployments, your replicaSets that were created from your YAML files, and your three nodes. You have now successfully created a Kubernetes cluster with one control-plane node and two worker nodes that are running fourteen pods.

Clean Up

Clean up is optional but is a good practice to save space on your local computer. To clean up what you just created, we need to delete the deployments and the nodes. You can delete both deployments in a single command:

```
kubectl delete deployments <httpd_deployment_name> <redis_deployment_name>
```

```
[Brandis-MBP:W17complex brandi$ kubectl delete deployments httpd-deploy redis-deploy  
deployment.apps "httpd-deploy" deleted  
deployment.apps "redis-deploy" deleted
```

To delete your cluster nodes, use the following command:

```
kubectl delete nodes <node_name> <node_name> <node_name>
```

```
[Brandis-MBP:W17complex brandi$ kubectl delete nodes k8cluster k8cluster-m02 k8cluster-m03  
node "k8cluster" deleted  
node "k8cluster-m02" deleted  
node "k8cluster-m03" deleted
```

To stop minikube, run the following command:

```
minikube stop
```

To delete minikube, run the following command:

```
minikube delete
```